

Designing for Failure

@italolelis

Desining **systems** for the **unexpected!**

Think about an airplane, a boeing 777 to be especific



Go Days Berlin 2019



(Electrical Load Management System)

Essentials

Resilience is a Requirement, Not a Feature



Liang Guo

Dependency Isolation and Graceful Degradation

Health-check and Load Balancing




```

func main() {
    // Create a new health instance
    h := health.New()

    sqlCheck, err := checkers.NewSQL(&checkers.SQLConfig{
        Pinger: db,
    })
    if err != nil {
        log.Fatalf("could not setup sql checker: %s", err)
    }

    if err = h.AddChecks([]*health.Config{
        {
            Name:      "db-reception-check",
            Checker:    sqlCheck,
            Interval:   time.Duration(3) * time.Second,
            Fatal:      true,
        },
        {
            Name:      "amqp-reception-check",
            Checker:    stream.NewChecker(stream.WithDSN(amqpDSN)),
            Interval:   time.Duration(3) * time.Second,
            Fatal:      true,
        },
    }); err != nil {
        log.Fatalf("could not add checkers: %s", err)
    }

    if err := h.Start(); err != nil {
        log.Fatalf("could not start health check: %s", err)
    }
}

```

If everything is OK you get...

```
HTTP/1.1 200 OK
Content-Length: 401
Content-Type: application/json
Date: Wed, 16 Jan 2019 19:59:49 GMT
```

```
{
  "details": {
    "amqp-reception-check": {
      "check_time": "2019-01-16T20:59:48.928127+01:00",
      "fatal": true,
      "first_failure_at": "0001-01-01T00:00:00Z",
      "name": "amqp-reception-check",
      "num_failures": 0,
      "status": "ok"
    },
    "db-reception-check": {
      "check_time": "2019-01-16T20:59:48.894341+01:00",
      "fatal": true,
      "first_failure_at": "0001-01-01T00:00:00Z",
      "name": "db-reception-check",
      "num_failures": 0,
      "status": "ok"
    }
  },
  "status": "ok"
}
```

If things are not good but your app
still can work...

```
HTTP/1.1 200 OK
Content-Length: 507
Content-Type: application/json
Date: Wed, 16 Jan 2019 20:03:28 GMT
```

```
{
  "details": {
    "amqp-reception-check": {
      "check_time": "2019-01-16T21:03:27.700167+01:00",
      "error":
"rabbitMQ health check failed on dial phase: dial tcp [::1]:5672: connect: connection refused",
      "first_failure_at": "2019-01-16T21:03:24.702708+01:00",
      "name": "amqp-reception-check",
      "num_failures": 2,
      "status": "failed"
    },
    "db-reception-check": {
      "check_time": "2019-01-16T21:03:27.698874+01:00",
      "fatal": true,
      "first_failure_at": "0001-01-01T00:00:00Z",
      "name": "db-reception-check",
      "num_failures": 0,
      "status": "ok"
    }
  },
  "status": "ok"
}
```

Otherwise...

HTTP/1.1 500 Internal Server Error
Content-Length: 588
Content-Type: application/json
Date: Wed, 16 Jan 2019 20:06:03 GMT

```
{
  "details": {
    "amqp-reception-check": {
      "check_time": "2019-01-16T21:06:03.702271+01:00",
      "error":
"rabbitMQ health check failed on dial phase: dial tcp [::1]:5672: connect: connection refused",
      "first_failure_at": "2019-01-16T21:03:24.702708+01:00",
      "name": "amqp-reception-check",
      "num_failures": 54,
      "status": "failed"
    },
    "db-reception-check": {
      "check_time": "2019-01-16T21:06:03.702259+01:00",
      "error": "dial tcp [::1]:5432: connect: connectionrefused",
      "fatal": true,
      "first_failure_at": "2019-01-16T21:05:12.706956+01:00",
      "name": "db-reception-check",
      "num_failures": 18,
      "status": "failed"
    }
  },
  "status": "failed"
}
```

Self-healing

Container orchestrators do this automatically

Load shedding

Circuit Breakers

```
func main() {  
    // Create a new fallback for when a circuit opens  
    fallbackFn := func(err error) error {  
        _, err := http.Post("post_to_channel_two")  
        return err  
    }  
  
    // Create a new hystrix-wrapped HTTP client  
    client := hystrix.NewClient(  
        hystrix.WithHTTPTimeout(200 * time.Millisecond),  
        hystrix.WithCommandName("MyCommand"),  
        hystrix.WithErrorPercentThreshold(20),  
        hystrix.WithSleepWindow(10),  
        hystrix.WithRequestVolumeThreshold(10),  
        hystrix.WithFallbackFunc(fallbackFn),  
    })  
    // Create an http.Request instance  
    req, _ := http.NewRequest(http.MethodGet, "http://google.com", nil)  
  
    // Call the `Do` method, which has a similar interface to the `http.Do` method  
    res, err := client.Do(req)  
    if err != nil { panic(err) }  
}
```


Circuit Closed

Hosts	581	90th	17ms
Median	6ms	99th	64ms
Mean	10ms	99.5th	114ms

CryptexDecipher

74,718 | 92 | 0.0 %
0 | 0 | 58

Host: 12.9/s

Cluster: 7,486.8/s

Circuit Closed

Hosts	581	90th	11ms
Median	3ms	99th	67ms
Mean	10ms	99.5th	405ms

VideoMetadataGetEpisodes

45,663 | 0 | 0.0 %
0 | 0 | 0

Host: 7.9/s

Cluster: 4,566.3/s

Circuit Closed

Hosts	581	90th	0ms
Median	0ms	99th	0ms
Mean	0ms	99.5th	0ms

GeoLookupCommand

27,455 | 0 | 0.0 %
0 | 0 | 9

Circuit Closed

Hosts	581	90th	1ms
Median	0ms	99th	34ms
Mean	1ms	99.5th	49ms

CinematchGetMovieRatings

60,672 | 2 | 0.0 %
0 | 0 | 0

Host: 10.4/s

Cluster: 6,067.4/s

Circuit Closed

Hosts	581	90th	41ms
Median	14ms	99th	143ms
Mean	21ms	99.5th	215ms

GPSGetGroup

41,142 | 0 | 0.0 %
0 | 0 | 0

Host: 7.1/s

Cluster: 4,114.2/s

Circuit Closed

Hosts	581	90th	114ms
Median	10ms	99th	318ms
Mean	38ms	99.5th	412ms

ABTestGetCellsAsAllocationsList

8,535 | 0 | 0.0 %
0 | 0 | 0

Circuit Closed

Hosts	581	90th	22ms
Median	3ms	99th	122ms
Mean	11ms	99.5th	312ms

VideoHistoryGetBookmarks

48,281 | 60 | 0.1 %
0 | 0 | 0

Host: 8.3/s

Cluster: 4,834.1/s

Circuit Closed

Hosts	581	90th	26ms
Median	8ms	99th	104ms
Mean	13ms	99.5th	158ms

SimpleDBGetItem

39,662 | 2 | 0.0 %
0 | 0 | 0

Host: 6.8/s

Cluster: 3,966.4/s

Circuit Closed

Hosts	581	90th	26ms
Median	14ms	99th	70ms
Mean	17ms	99.5th	123ms

ABTestGetAllocationsCommand

7,422 | 0 | 0.0 %
0 | 0 | 0

Retry Logic

```

func main() {
    // Exponential Backoff increases the backoff at an exponential rate
    initTimeout := 2*time.Millisecond
    maxTimeout := 10*time.Millisecond
    expFactor := 2
    maxJitterInterval := 2*time.Millisecond

    backoff := heimdall.NewExponentialBackoff(
        initTimeout,
        maxTimeout,
        expFactor,
        maxJitterInterval,
    )

    // Create a new retry mechanism with the backoff
    retrier := heimdall.NewRetrier(backoff)

    // Create an http client with the retry mechanism, and the number of times you would like to retry
    client := httpclient.NewClient(
        httpclient.WithHTTPTimeout(1000 * time.Millisecond),
        httpclient.WithRetrier(retrier),
        httpclient.WithRetryCount(4),
    )

    // Create an http.Request instance
    req, _ := http.NewRequest(http.MethodGet, "http://google.com", nil)

    // Call the `Do` method, which has a similar interface to the `http.Do` method
    res, err := client.Do(req)
    if err != nil { panic(err) }
}

```

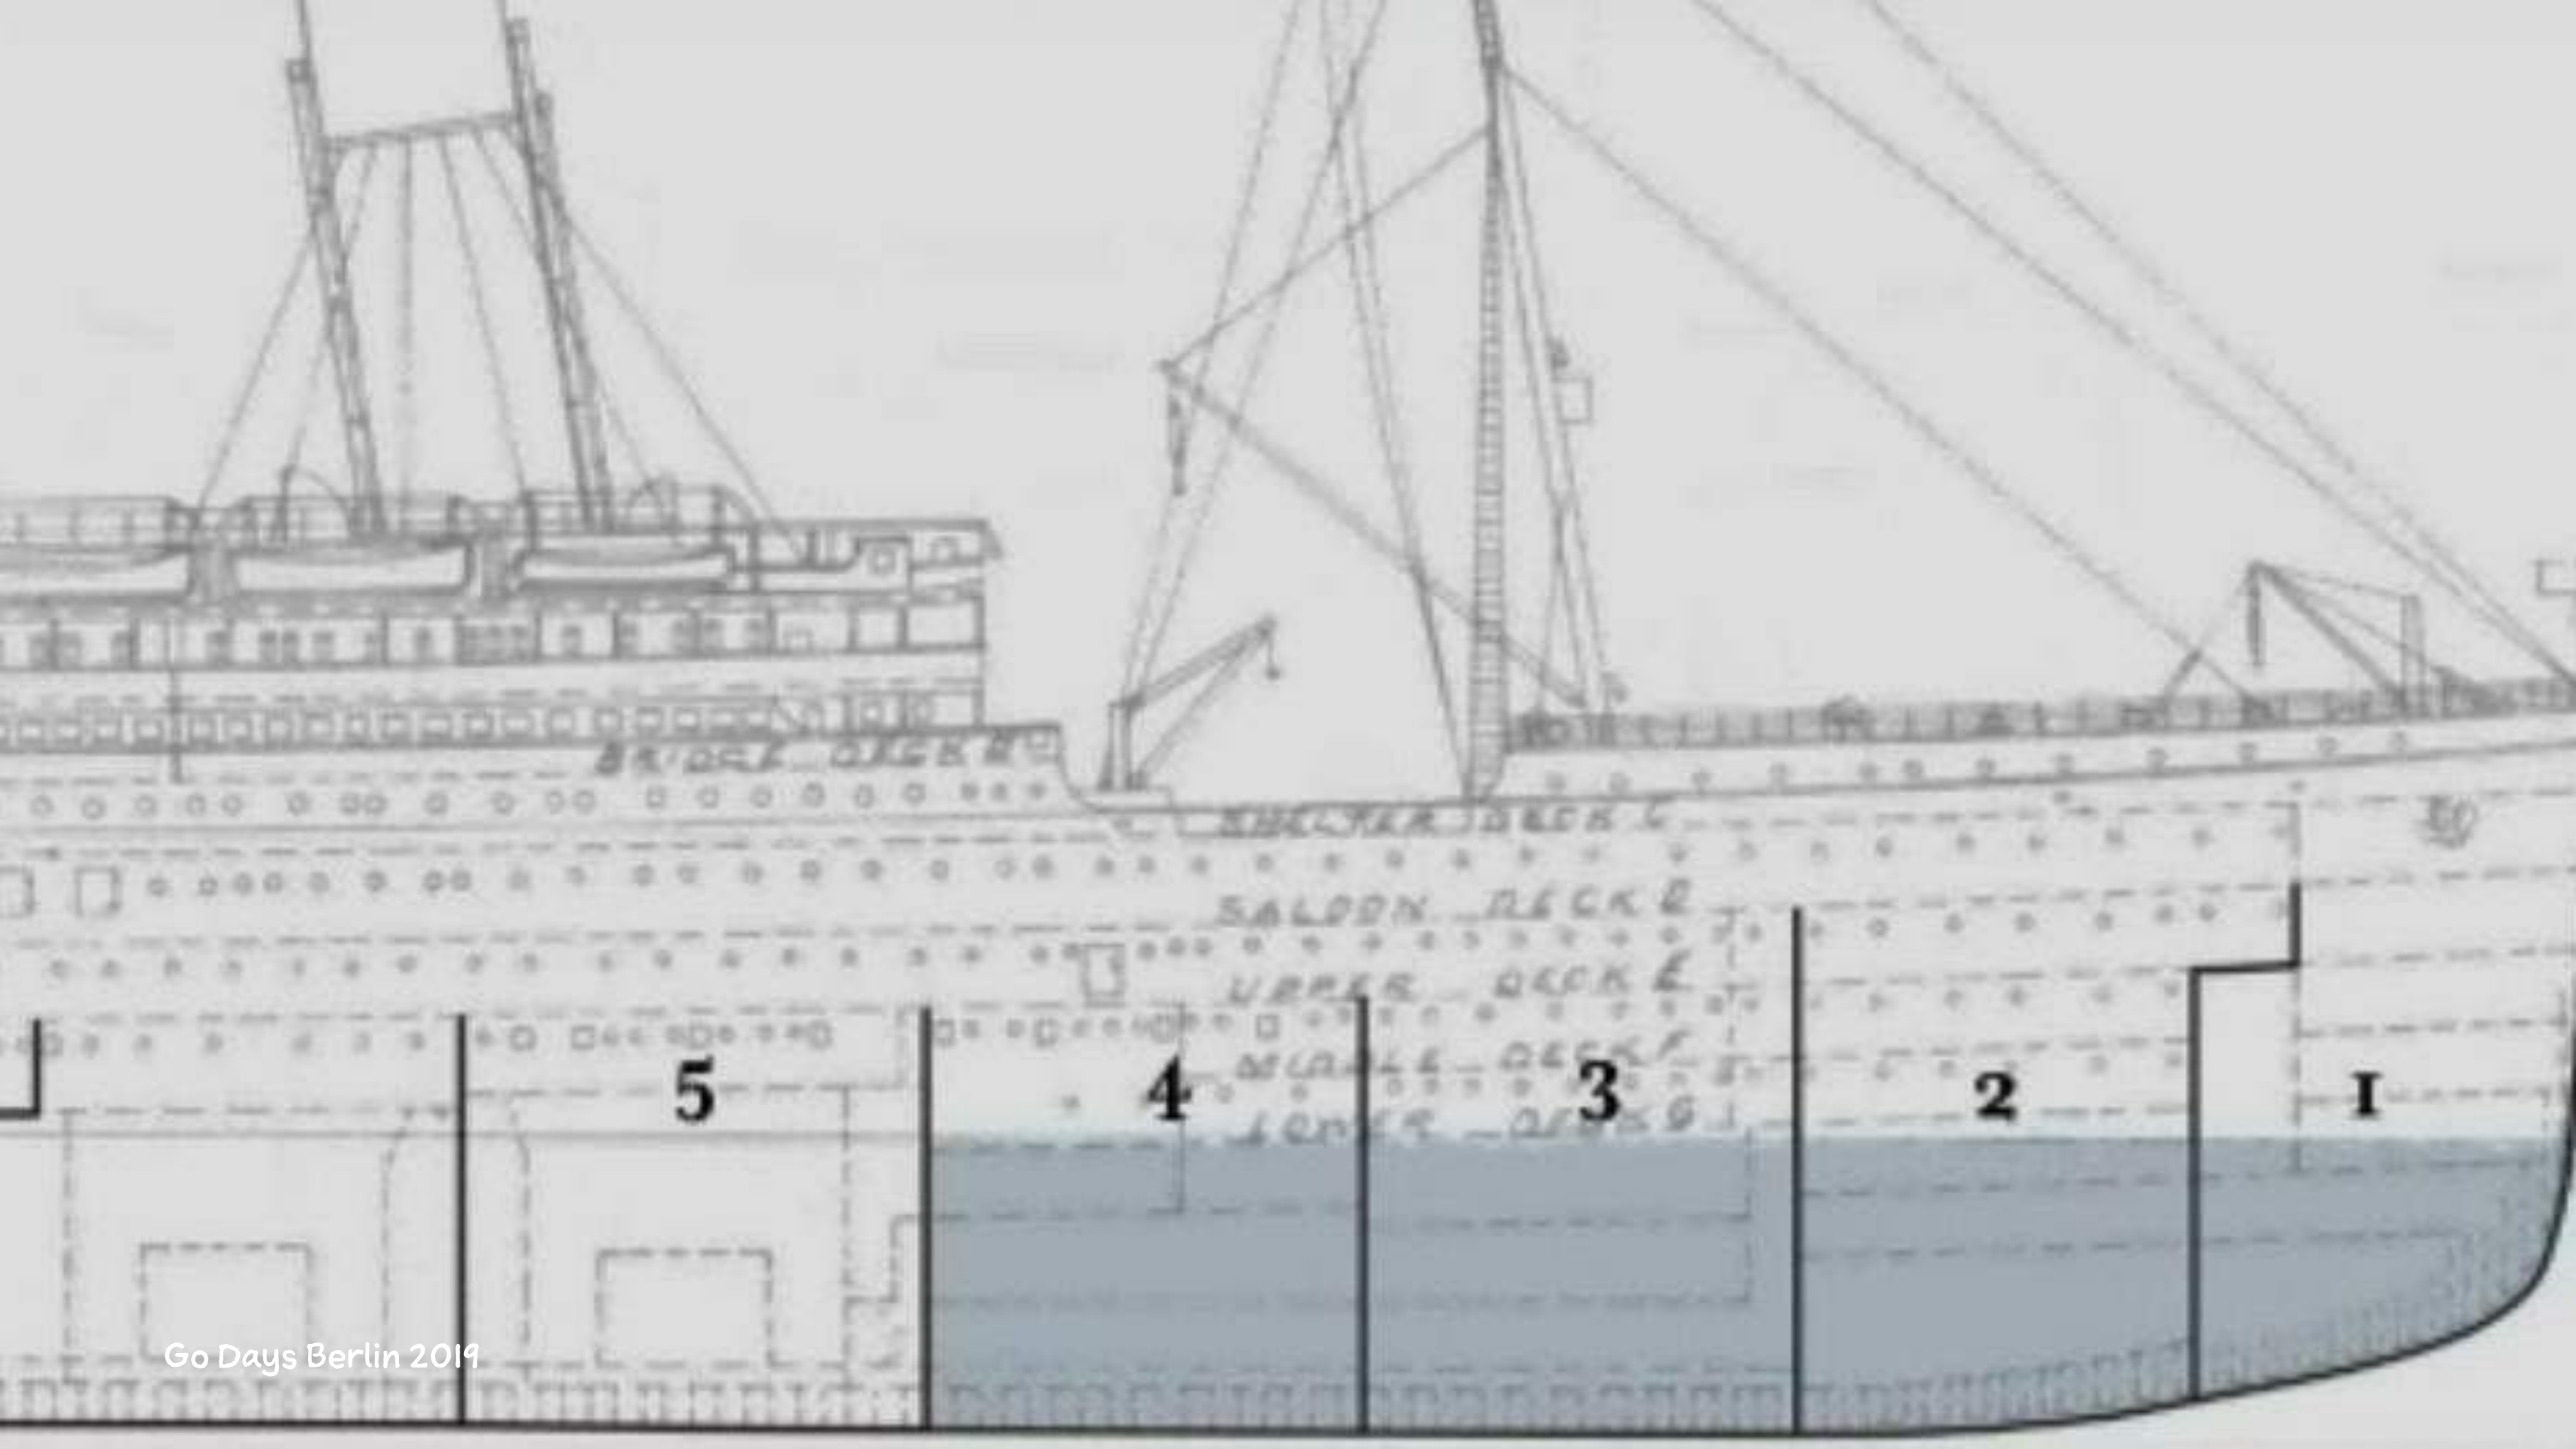
Rate Limiters


```
func main() {  
    rate, err := limiter.NewRateFromFormatted("1000-H")  
    if err != nil {  
        panic(err)  
    }  
  
    store := memory.NewStore()  
  
    // Then, create the limiter instance which takes the store and the rate as arguments.  
    // Now, you can give this instance to any supported middleware.  
    instance := limiter.New(store, rate)  
}
```

Bulkhead



Go Days Berlin 2019



Outlier Server Host Detection

Outbox Pattern

```

func main() {
    // ...

    ds, err := postgres.WithInstance(ctx, db.DB)
    if err != nil {
        log.Fatalf("could not create a postgres wrapper", err)
    }

    o, err := outboxer.New(
        outboxer.WithDataStore(ds),
        outboxer.WithEventStream(amqp.NewAMQP(conn)),
        outboxer.WithCheckInterval(1*time.Second),
        outboxer.WithCleanupInterval(5*time.Second),
    )
    defer o.Stop()

    // Start the listeners for sending and cleaning messages
    o.Start()

    // Sends a message
    if err = o.Send(ctx, &outboxer.OutboxMessage{
        Payload: []byte("test payload"),
        Options: map[string]interface{}{
            amqp.ExchangeNameOption: "test",
            amqp.ExchangeTypeOption: "test.send",
        },
    }); err != nil {
        log.Fatalf("could not send message: %s", err)
    }
}

```

Service Mesh

Observability

Defining your SLO's and SLI's



Monitoring

```
if err := view.Register(  
    ochttp.ClientSentBytesDistribution,  
    ochttp.ClientReceivedBytesDistribution,  
    ochttp.ClientRoundtripLatencyDistribution,  
); err != nil {  
    logger.Fatal(err)  
}  
  
exporter, err := prometheus.NewExporter(prometheus.Options{  
    Namespace: cfg.ServiceName,  
})  
if err != nil {  
    log.Fatal("failed to create the prometheus stats exporter")  
}  
view.RegisterExporter(exporter)  
view.SetReportingPeriod(cfg.ReportingPeriod)
```

Distributed Tracing

```
exporter, err := jaeger.NewExporter(jaeger.Options{
    CollectorEndpoint: cfg.CollectorEndpoint,
    Process: jaeger.Process{
        ServiceName: cfg.ServiceName,
    },
})
if err != nil {
    log.Error("could not create the jaeger exporter")
}

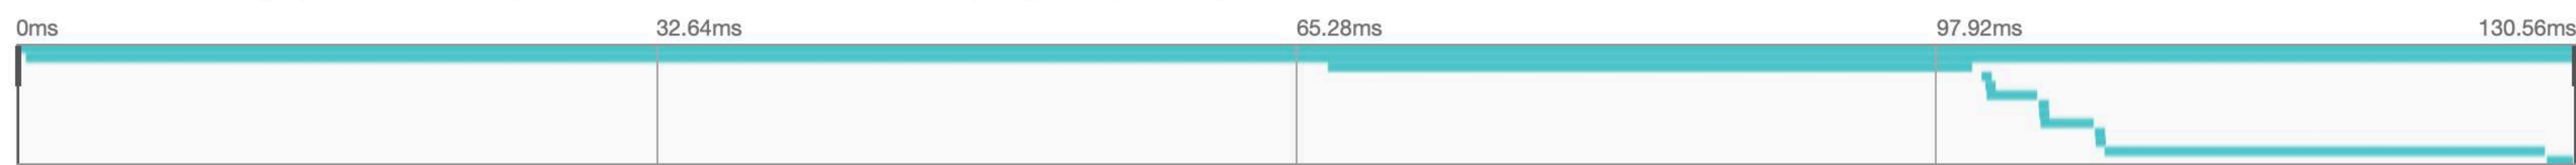
trace.RegisterExporter(exporter)
trace.ApplyConfig(trace.Config{DefaultSampler: trace.AlwaysSample()})
```


▼ reception: Recv./orders

Search...

View Options ▼

Trace Start: **January 7, 2019 9:25 PM** | Duration: **130.56ms** | Services: **1** | Depth: **3** | Total Spans: **13**



Service & Operation	0ms	32.64ms	65.28ms	97.92ms	130.56ms
---------------------	-----	---------	---------	---------	----------

Sent.sql:query

Service: **reception** | Duration: **32.86ms** | Start Time: **66.9ms**

> **Tags:** sql.query = SELECT * FROM coffees WHERE name = \$1 LIMIT 1 | sql.arg.1 = cappuccino | status.code = 0 | status.message =

> **Process:**

SpanID: 7aca3319934124e7

reception	Sent.sql:rows_next				0.07ms	
reception	Sent.sql:rows_close				0.01ms	
reception	Sent.sql:query				2.6ms	<div></div>
reception	Sent.sql:rows_next				0.01ms	
reception	Sent.sql:rows_close				0.03ms	
reception	Sent.sql:query				2.72ms	<div></div>
reception	Sent.sql:rows_next				0.01ms	
reception	Sent.sql:rows_close				0.01ms	
reception	Sent.sql:exec				22.47ms	<div></div>

Open Census

```
import (  
    "go.opencensus.io/exporter/prometheus"  
    "go.opencensus.io/plugin/ochttp"  
    "go.opencensus.io/stats/view"  
)
```

Recap

1. Always think about your dependencies
2. Dependency Isolation and Graceful Degradation
3. Load shedding and Request Controlling
4. Observability is not optional

Questions and links!

- Example application: <https://github.com/italolelis/coffee-shop>
- Link to the slides: <https://github.com/italolelis/talks>



Thank you!