


Designing for Failure

@italolelis


Think about an airplane, a boeing 777 to be especific



Go Days Berlin 2019



ELMS



Essentials

Resilience is a Requirement, Not a Feature



Liang Guo

Dependency Isolation and Graceful Degradation

Health-check and Load Balancing



```

import (
    "net/http"
    "time"

    "github.com/hellofresh/health-go"
    healthMySQL "github.com/hellofresh/health-go/checks/mysql"
)

func main() {
    health.Register(health.Config{
        Name: "kafka",
        Timeout: time.Second*5,
        SkipOnError: true,
        Check: func() error {
            // kafka health check implementation goes here
        },
    })

    health.Register(health.Config{
        Name: "mysql",
        Timeout: time.Second * 2,
        SkipOnError: false,
        Check: healthMySQL.New(healthMySQL.Config{
            DSN: "test:test@tcp(0.0.0.0:31726)/test?charset=utf8",
        },
    })

    http.Handle("/status", health.Handler())
    http.ListenAndServe(":8080", nil)
}

```


If everything is OK you get...

```
{
  "status": "OK",
  "timestamp": "2017-01-01T00:00:00.413567856+033:00",
  "system": {
    "version": "go1.8",
    "goroutines_count": 4,
    "total_alloc_bytes": 21321,
    "heap_objects_count": 21323,
    "alloc_bytes": 234523
  }
}
```

If things are not good but your app
still can work...

```
{
  "status": "Partially Available",
  "timestamp": "2017-01-01T00:00:00.413567856+033:00",
  "failures": {
    "rabbitmq": "Failed during rabbitmq health check"
  },
  "system": {
    "version": "go1.8",
    "goroutines_count": 4,
    "total_alloc_bytes": 21321,
    "heap_objects_count": 21323,
    "alloc_bytes": 234523
  }
}
```

Otherwise...

```
{
  "status": "Unavailable",
  "timestamp": "2017-01-01T00:00:00.413567856+033:00",
  "failures": {
    "mongodb": "Failed during mongodb health check"
  },
  "system": {
    "version": "go1.8",
    "goroutines_count": 4,
    "total_alloc_bytes": 21321,
    "heap_objects_count": 21323,
    "alloc_bytes": 234523
  }
}
```


Self-healing

In **kube** this is as simple as defining a
YAML file rule

```
apiVersion: apps/v1beta1
kind: Deployment
metadata:
  name: coffee-shop-deploy
spec:
  replicas: 2
  template:
    metadata:
      labels:
        app: coffee-shop
    spec:
      containers:
        - name: coffee-shop
          image: italoelalis/coffee-shop:0.5.0
          ports:
            - containerPort: 9876
          env:
            - name: VERSION
              value: "0.9"
```

Load shedding

Circuit Breakers

```
func main() {  
    // Create a new fallback for when a circuit opens  
    fallbackFn := func(err error) error {  
        _, err := http.Post("post_to_channel_two")  
        return err  
    }  
  
    // Create a new hystrix-wrapped HTTP client  
    client := hystrix.NewClient(  
        hystrix.WithHTTPTimeout(200 * time.Millisecond),  
        hystrix.WithCommandName("MyCommand"),  
        hystrix.WithErrorPercentThreshold(20),  
        hystrix.WithSleepWindow(10),  
        hystrix.WithRequestVolumeThreshold(10),  
        hystrix.WithFallbackFunc(fallbackFn),  
    })  
    // Create an http.Request instance  
    req, _ := http.NewRequest(http.MethodGet, "http://google.com", nil)  
  
    // Call the `Do` method, which has a similar interface to the `http.Do` method  
    res, err := client.Do(req)  
    if err != nil { panic(err) }  
}
```


Retry Logic

```

func main() {
    // Exponential Backoff increases the backoff at an exponential rate
    initTimeout := 2*time.Millisecond
    maxTimeout := 10*time.Millisecond
    expFactor := 2
    maxJitterInterval := 2*time.Millisecond

    backoff := heimdall.NewExponentialBackoff(
        initTimeout,
        maxTimeout,
        expFactor,
        maxJitterInterval,
    )

    // Create a new retry mechanism with the backoff
    retrier := heimdall.NewRetrier(backoff)

    // Create an http client with the retry mechanism, and the number of times you would like to retry
    client := httpclient.NewClient(
        httpclient.WithHTTPTimeout(1000 * time.Millisecond),
        httpclient.WithRetrier(retrier),
        httpclient.WithRetryCount(4),
    )

    // Create an http.Request instance
    req, _ := http.NewRequest(http.MethodGet, "http://google.com", nil)

    // Call the `Do` method, which has a similar interface to the `http.Do` method
    res, err := client.Do(req)
    if err != nil { panic(err) }
}

```

Bulkhead

Rate Limiters

```
func main() {  
    rate, err := limiter.NewRateFromFormatted("1000-H")  
    if err != nil {  
        panic(err)  
    }  
  
    store := memory.NewStore()  
  
    // Then, create the limiter instance which takes the store and the rate as arguments.  
    // Now, you can give this instance to any supported middleware.  
    instance := limiter.New(store, rate)  
}
```

Outbox Pattern

Outlier Server Host Detection

Service Mesh

SLO's and SLI's



Monitoring

```
if err := view.Register(  
    ochttp.ClientSentBytesDistribution,  
    ochttp.ClientReceivedBytesDistribution,  
    ochttp.ClientRoundtripLatencyDistribution,  
); err != nil {  
    logger.Fatal(err)  
}  
  
exporter, err := prometheus.NewExporter(prometheus.Options{  
    Namespace: cfg.ServiceName,  
})  
if err != nil {  
    log.Fatal("failed to create the prometheus stats exporter")  
}  
view.RegisterExporter(exporter)  
view.SetReportingPeriod(cfg.ReportingPeriod)
```

Distributed Tracing


```
exporter, err := jaeger.NewExporter(jaeger.Options{
    CollectorEndpoint: cfg.CollectorEndpoint,
    Process: jaeger.Process{
        ServiceName: cfg.ServiceName,
    },
})
if err != nil {
    log.Error("could not create the jaeger exporter")
}


trace.RegisterExporter(exporter)
trace.ApplyConfig(trace.Config{DefaultSampler: trace.AlwaysSample()})
```

▼ reception: Recv./orders

Search...

View Options ▼

Trace Start: **January 7, 2019 9:25 PM** | Duration: **130.56ms** | Services: **1** | Depth: **3** | Total Spans: **13**



Service & Operation	0ms	32.64ms	65.28ms	97.92ms	130.56ms
---------------------	-----	---------	---------	---------	----------

Sent.sql:query

Service: **reception** | Duration: **32.86ms** | Start Time: **66.9ms**

> **Tags:** sql.query = SELECT * FROM coffees WHERE name = \$1 LIMIT 1 | sql.arg.1 = cappuccino | status.code = 0 | status.message =

> **Process:**

SpanID: 7aca3319934124e7

reception	Sent.sql:rows_next				0.07ms	
reception	Sent.sql:rows_close				0.01ms	
reception	Sent.sql:query				2.6ms	<div></div>
reception	Sent.sql:rows_next				0.01ms	
reception	Sent.sql:rows_close				0.03ms	
reception	Sent.sql:query				2.72ms	<div></div>
reception	Sent.sql:rows_next				0.01ms	
reception	Sent.sql:rows_close				0.01ms	
reception	Sent.sql:exec				22.47ms	<div></div>

Open Census

```
import (  
    "go.opencensus.io/exporter/prometheus"  
    "go.opencensus.io/plugin/ochttp"  
    "go.opencensus.io/stats/view"  
)
```

Recap

1. Always think about your dependencies
2. Dependency Isolation and Graceful Degradation
3. Load shedding and Request Controlling
4. Observability is not optional

Questions and links!

- Example application: <https://github.com/italolelis/coffee-shop>
- Link to the slides: <https://github.com/italolelis/talks>



Thank you!